

ML_Prac3

October 18, 2023

0.0.1 Aim: Given a bank customer, build a neural network-based classifier that can determine whether they will leave or not in the next 6 months. **Dataset Description:** The case study is from an open-source dataset from Kaggle. The dataset contains 10,000 sample points with 14 distinct features such as CustomerId, CreditScore, Geography, Gender, Age, Tenure, Balance, etc.

Name: Chinmay Gokhale

Div: BE-A

Roll No. B211047

0.0.2 Import the Libraries

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[ ]:
```

```
[2]: df=pd.read_csv("Churn_Modelling.csv")
```

```
[3]: df
```

```
[3]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	
3	4	15701354	Boni	699	France	Female	39	
4	5	15737888	Mitchell	850	Spain	Female	43	
...	
9995	9996	15606229	Obijiaku	771	France	Male	39	
9996	9997	15569892	Johnstone	516	France	Male	35	
9997	9998	15584532	Liu	709	France	Female	36	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	
9999	10000	15628319	Walker	792	France	Female	28	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	
...	
9995	5	0.00	2	1	0	
9996	10	57369.61	1	1	1	
9997	7	0.00	1	0	1	
9998	3	75075.31	2	1	0	
9999	4	130142.79	1	1	0	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0
...
9995	96270.64	0
9996	101699.77	0
9997	42085.58	1
9998	92888.52	1
9999	38190.78	0

[10000 rows x 14 columns]

```
[4]: df.head()
```

```
[4]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	
3	4	15701354	Boni	699	France	Female	39	
4	5	15737888	Mitchell	850	Spain	Female	43	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0

2	113931.57	1
3	93826.63	0
4	79084.10	0

```
[5]: df.tail()
```

```
[5]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
9995	9996	15606229	Obijiaku	771	France	Male	39	
9996	9997	15569892	Johnstone	516	France	Male	35	
9997	9998	15584532	Liu	709	France	Female	36	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	
9999	10000	15628319	Walker	792	France	Female	28	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
9995	5	0.00	2	1	0	
9996	10	57369.61	1	1	1	
9997	7	0.00	1	0	1	
9998	3	75075.31	2	1	0	
9999	4	130142.79	1	1	0	

	EstimatedSalary	Exited
9995	96270.64	0
9996	101699.77	0
9997	42085.58	1
9998	92888.52	1
9999	38190.78	0

```
[6]: df.shape
```

```
[6]: (10000, 14)
```

```
[7]: df.describe()
```

```
[7]:
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	\
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	

	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
count	10000.000000	10000.000000	10000.000000	10000.000000	
mean	76485.889288	1.530200	0.70550	0.515100	
std	62397.405202	0.581654	0.45584	0.499797	

min	0.000000	1.000000	0.00000	0.000000
25%	0.000000	1.000000	0.00000	0.000000
50%	97198.540000	1.000000	1.00000	1.000000
75%	127644.240000	2.000000	1.00000	1.000000
max	250898.090000	4.000000	1.00000	1.000000

	EstimatedSalary	Exited
count	10000.000000	10000.000000
mean	100090.239881	0.203700
std	57510.492818	0.402769
min	11.580000	0.000000
25%	51002.110000	0.000000
50%	100193.915000	0.000000
75%	149388.247500	0.000000
max	199992.480000	1.000000

```
[8]: df.isnull()
```

```
[8]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	
...	
9995	False	False	False	False	False	False	False	
9996	False	False	False	False	False	False	False	
9997	False	False	False	False	False	False	False	
9998	False	False	False	False	False	False	False	
9999	False	False	False	False	False	False	False	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	False	False	False	False	False	
1	False	False	False	False	False	
2	False	False	False	False	False	
3	False	False	False	False	False	
4	False	False	False	False	False	
...	
9995	False	False	False	False	False	
9996	False	False	False	False	False	
9997	False	False	False	False	False	
9998	False	False	False	False	False	
9999	False	False	False	False	False	

	EstimatedSalary	Exited
0	False	False
1	False	False

```

2           False  False
3           False  False
4           False  False
...
9995        False  False
9996        False  False
9997        False  False
9998        False  False
9999        False  False

```

[10000 rows x 14 columns]

```
[9]: df.isnull().sum()
```

```

[9]: RowNumber      0
    CustomerId      0
    Surname         0
    CreditScore     0
    Geography       0
    Gender          0
    Age            0
    Tenure          0
    Balance         0
    NumOfProducts  0
    HasCrCard       0
    IsActiveMember  0
    EstimatedSalary 0
    Exited          0
    dtype: int64

```

```
[10]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   RowNumber             10000 non-null  int64
 1   CustomerId            10000 non-null  int64
 2   Surname               10000 non-null  object
 3   CreditScore           10000 non-null  int64
 4   Geography             10000 non-null  object
 5   Gender                10000 non-null  object
 6   Age                   10000 non-null  int64
 7   Tenure                10000 non-null  int64
 8   Balance                10000 non-null  float64
 9   NumOfProducts         10000 non-null  int64

```

```

10 HasCrCard      10000 non-null  int64
11 IsActiveMember 10000 non-null  int64
12 EstimatedSalary 10000 non-null  float64
13 Exited         10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB

```

```
[11]: df.info
```

```

[11]: <bound method DataFrame.info of
CreditScore Geography Gender Age \
0          1    15634602   Hargrave      619    France  Female   42
1          2    15647311     Hill      608     Spain  Female   41
2          3    15619304     Onio      502    France  Female   42
3          4    15701354     Boni      699    France  Female   39
4          5    15737888  Mitchell      850     Spain  Female   43
...
9995      9996    15606229   Obijiaku      771    France   Male   39
9996      9997    15569892  Johnstone      516    France   Male   35
9997      9998    15584532     Liu      709    France  Female   36
9998      9999    15682355  Sabbatini      772   Germany   Male   42
9999     10000    15628319    Walker      792    France  Female   28

Tenure  Balance  NumOfProducts  HasCrCard  IsActiveMember \
0        2      0.00             1          1              1
1        1  83807.86             1          0              1
2        8 159660.80             3          1              0
3        1      0.00             2          0              0
4        2 125510.82             1          1              1
...
9995      5      0.00             2          1              0
9996     10  57369.61             1          1              1
9997      7      0.00             1          0              1
9998      3  75075.31             2          1              0
9999      4 130142.79             1          1              0

EstimatedSalary  Exited
0          101348.88      1
1          112542.58      0
2          113931.57      1
3           93826.63      0
4           79084.10      0
...
9995          96270.64      0
9996         101699.77      0
9997          42085.58      1
9998          92888.52      1

```

```
9999          38190.78          0
```

```
[10000 rows x 14 columns]>
```

```
[12]: df.columns
```

```
[12]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',  
          'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',  
          'IsActiveMember', 'EstimatedSalary', 'Exited'],  
          dtype='object')
```

```
[13]: df.dtypes
```

```
[13]: RowNumber          int64  
      CustomerId       int64  
      Surname         object  
      CreditScore      int64  
      Geography       object  
      Gender          object  
      Age             int64  
      Tenure          int64  
      Balance         float64  
      NumOfProducts   int64  
      HasCrCard       int64  
      IsActiveMember  int64  
      EstimatedSalary float64  
      Exited          int64  
      dtype: object
```

0.0.3 Splitting the data

```
[14]: x = df.  
      ↪drop(["RowNumber","CustomerId","Surname","Geography","Gender","Exited"],axis=1)
```

```
[15]: x
```

```
[15]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	\
0	619	42	2	0.00	1	1	
1	608	41	1	83807.86	1	0	
2	502	42	8	159660.80	3	1	
3	699	39	1	0.00	2	0	
4	850	43	2	125510.82	1	1	
...	
9995	771	39	5	0.00	2	1	
9996	516	35	10	57369.61	1	1	
9997	709	36	7	0.00	1	0	
9998	772	42	3	75075.31	2	1	

9999	792	28	4	130142.79	1	1
------	-----	----	---	-----------	---	---

	IsActiveMember	EstimatedSalary
0	1	101348.88
1	1	112542.58
2	0	113931.57
3	0	93826.63
4	1	79084.10
...
9995	0	96270.64
9996	1	101699.77
9997	1	42085.58
9998	0	92888.52
9999	0	38190.78

[10000 rows x 8 columns]

```
[16]: y=df["Exited"]
```

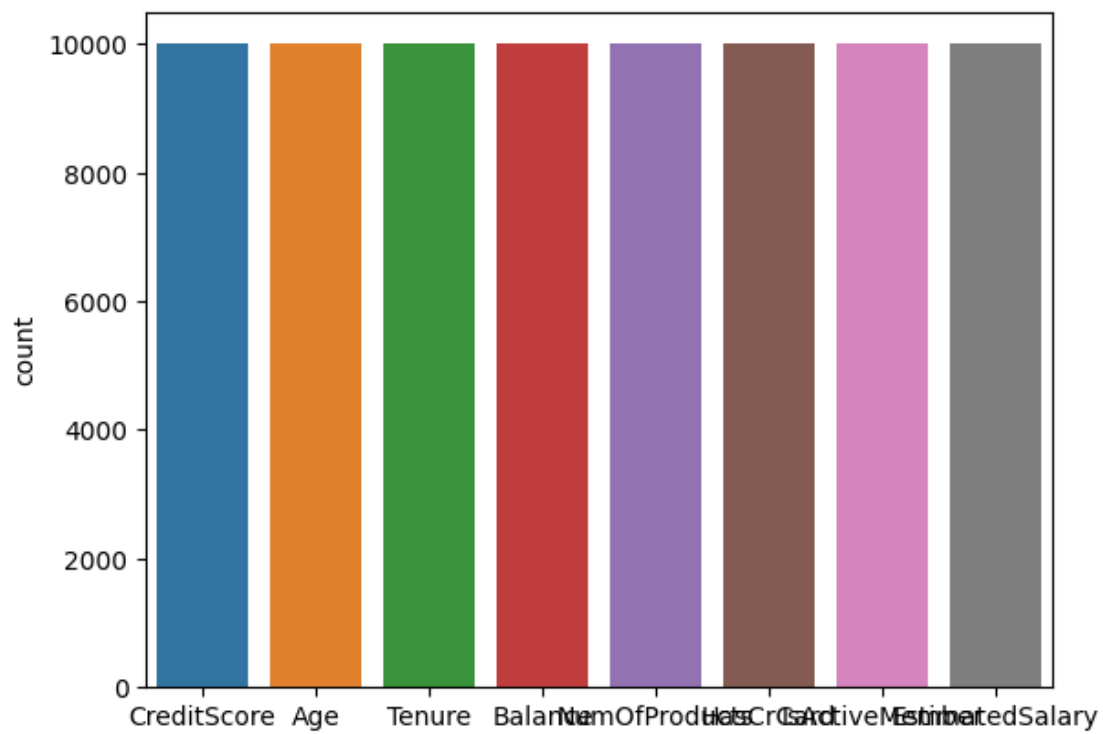
```
[17]: y
```

```
[17]: 0      1
      1      0
      2      1
      3      0
      4      0
      ..
      9995    0
      9996    0
      9997    1
      9998    1
      9999    0
      Name: Exited, Length: 10000, dtype: int64
```

1 Checking Balancing of data


```
[18]: sns.countplot(x)
```

```
[18]: <Axes: ylabel='count'>
```


```
[19]: sns.countplot(y)
```

```
[19]: <Axes: ylabel='count'>
```



```
[20]: sns.countplot(x=y)
```

```
[20]: <Axes: xlabel='Exited', ylabel='count'>
```



```
[21]: y.value_counts()
```

```
[21]: Exited
0    7963
1     2037
Name: count, dtype: int64
```

```
[22]: x.value_counts()
```

```
[22]: CreditScore  Age  Tenure  Balance  NumOfProducts  HasCrCard  IsActiveMember
EstimatedSalary
350           39    0      109733.20  2              0              0
123602.11           1
695           34    9        0.00    2              1              1
67502.12           1
28           5      171069.39  2              1              1
88689.40           1
29           5        0.00    2              1              1
6770.44           1
9            0.00    2              1              0
111565.45          1
```

```

..
608      33  9      89968.69  1      1      0
68777.26      1
      34  3      106288.54  1      1      1
36639.25      1
      4      88772.87  1      1      1
168822.01      1
      7      86656.13  1      0      1
59890.29      1
850      81  5      0.00      2      1      1
44827.47      1
Name: count, Length: 10000, dtype: int64

```

1.1 Feature scaling

```
[23]: from sklearn.preprocessing import StandardScaler
      sc=StandardScaler()
```

```
[24]: X_scale= sc.fit_transform(x)
```

```
[25]: X_scale
```

```
[25]: array([[ -0.32622142,  0.29351742, -1.04175968, ...,  0.64609167,
          0.97024255,  0.02188649],
        [ -0.44003595,  0.19816383, -1.38753759, ..., -1.54776799,
          0.97024255,  0.21653375],
        [ -1.53679418,  0.29351742,  1.03290776, ...,  0.64609167,
         -1.03067011,  0.2406869 ],
        ...,
        [  0.60498839, -0.27860412,  0.68712986, ..., -1.54776799,
          0.97024255, -1.00864308],
        [  1.25683526,  0.29351742, -0.69598177, ...,  0.64609167,
         -1.03067011, -0.12523071],
        [  1.46377078, -1.04143285, -0.35020386, ...,  0.64609167,
         -1.03067011, -1.07636976]])
```

2 Cross Validation

```
[26]: from sklearn.model_selection import train_test_split
      X_train,X_test,Y_train,Y_test=␣
      ↪train_test_split(X_scale,y,random_state=47,test_size=0.47)
```

```
[27]: X_train.shape
```

```
[27]: (5300, 8)
```

```
[28]: X_test.shape
```

```
[28]: (4700, 8)
```

2.1 Initialize and build the model.

```
[29]: from sklearn.neural_network import MLPClassifier
ann=MLPClassifier(hidden_layer_sizes=(100,100,100),random_state=2,activation="relu")
```

```
[30]: ann.fit(X_train, Y_train)
```


```
[30]: MLPClassifier(hidden_layer_sizes=(100, 100, 100), random_state=2)
```

```
[31]: y_pred = ann.predict(X_test)
y_pred
```

```
[31]: array([0, 0, 1, ..., 0, 0, 0])
```

```
[32]: from sklearn.metrics import ConfusionMatrixDisplay, accuracy_score, classification_report
ConfusionMatrixDisplay.from_predictions(Y_test,y_pred)
```

```
[32]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7f10943391b0>
```



```
[33]: accuracy_score(Y_test,y_pred)
```

```
[33]: 0.8034042553191489
```

```
[34]: print(classification_report(Y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.88	0.88	3755
1	0.51	0.50	0.51	945
accuracy			0.80	4700
macro avg	0.69	0.69	0.69	4700
weighted avg	0.80	0.80	0.80	4700

```
[ ]:
```